

Vantemio

AI automation core
Studio, services and commerce

Pre-seed proposal
8 October 2026

Vitalii Pavlichuk
Founder and Product Architect



Core and application modules

Reusable orchestration and execution core

One core links task context, tools, execution records and accepted output versions.

Vantemio Studio

First testing and monetization stage: editing, motion graphics and audio for Vantemio's own channels.

Vantemio Market

Planned marketplace for ready-made assets, custom services and independent sellers, using Base and Boson Protocol.

Next modules: Social, Marketing, CRM and Donations. Integrated production remains in development.

The production problem

Cloud tokens spent on routine work

Repeated prompts, corrections and context transfers increase cloud usage and production time.

Artistic production and reusable experience

Direction, pacing and sound need coherent execution. Recorded lessons should improve the next production.

Quality, rights and trustworthy sales

Buyers need clear terms, verified files and correct delivery. Crypto settlement still carries network and protocol fees.

Creators
Small studios
Agencies

Short videos can still require long production runs

12 hours

Reported by @anabology for an AI-directed video produced from a prompt, tools and a moodboard.

Minutes to hours

Tens of minutes for simple work, hours for complex graphics. This remains a target for equivalent tasks and quality.

Track manual prompts, corrections and elapsed production time. Measure rendering separately.

[Reference: @anabology's production example](#)

How production time should decrease

Plan before execution

Keep scene intent, timing, rights and sound requirements in structured plans shared by downstream tools.

Use available machines

Distribute independent admitted jobs and graphics work across bounded workers under Station control.


Repair only failed work

Preserve accepted results. Recheck failed components and their dependencies before assembling a new version.

Benchmark the complete job


Track elapsed time, render time, cost, output quality and human interventions on equivalent briefs.

The intended end-to-end service cycle



One recorded workflow carries task context, accepted assets and execution history between stages.

The reusable technical core



Bounded executors and service adapters perform the admitted work and return results.

Python core, Remotion and Blender graphics, FFmpeg and Vidra/MLT, Node.js and Solidity

Production components in development

Video and sound matrices

Scene timing, visual assets, voice, music and effects are represented in structured production plans.

Audio workflows

Voice, music and event sounds have separate timing and processing rules. Mixing and quality integration continue.

Editing and graphics

Source includes assembly, graphics, transitions and render logic. Reliable integrated output is still being developed.

Memory and repair

Memory, registry and repair mechanisms retain task context and recorded outcomes for later decisions.

Working today: automated YouTube competitor analysis. Full Studio execution remains the next integration target.

Video and audio are structured production

Video matrix

Scenes, shot intent, assets, timing, transitions and composition become explicit instructions for assembly and review.

Creative direction in code

Rules connect script intent with pacing, scene continuity, visual presentation and the audio plan.

Sound matrix

Narration, music and effects are separate layers with timing, synchronization, processing and mixing instructions.

Current engineering work

Improve output quality, remove integration failures and make complete productions repeatable.

Experience-based learning



Persistent context

The design retains task history, versions, errors and decisions in persistent memory and graph relationships.

Controlled improvement

Approved lessons guide direction, prompt design, editing, motion graphics and sound in later tasks.

Planned shared learning uses isolated customer workspaces. Model training needs a separate data-use basis.

Station controls distributed execution

Python orchestration core

Task admission, persistent memory, registries and production state remain under the central workstation.

Station

RTX 5070 Ti workstation
Owns queues and orchestration
Runs local computation

Windows laptop

Qwen Coder 7B helper
Receives admitted tasks
Returns results and receipts

MacBook

Development and client work
Receives bounded jobs
Returns outputs and receipts

The registry admits additional compute nodes. Readiness and resource budgets govern each assignment.

Cloud and local AI roles

Model / role	Assigned work	Integration requirement
Cloud reasoning	Direction, dramaturgy and complex planning	Approved access, bounded cost and validated output
Qwen 27B / Coder 7B	Local planning and bounded code tasks	Resource admission and recorded request results
Vision adapters	Scene and frame screening	Coverage, task context and separate quality review
Qwen3-TTS / Whisper / VAD	Narration and timestamped speech analysis	Voice rights, timing and intelligibility checks
ACE-Step / FFmpeg / CLAP	Music, sound processing and semantic analysis	Commercial model clearance and mix validation

Replaceable baseline. MMAudio and AudioLDM remain research components pending rights and quality checks.


Cloud reasoning and local execution

Fewer tokens

Cloud models handle global planning. Local models and tools carry out routine production tasks.

Measure tokens, retries, human work and cost per accepted output on the same brief.

External generation remains the initial path. Local generation is a planned, workload-tested option.



Models can change with hardware and workload. Adapters preserve task schemas, validation and coordinator authority.

Configured production workflows

Initial configuration

Define a brand profile, sources, channels, schedule, budget and authorization rules before running a workflow.

Autonomous operating target

The target loop researches, plans, produces and reviews content, then publishes through permitted interfaces.

Publication authorization and recovery

Bind publication to the approved account and asset version. Reconcile interrupted operations before retrying.

First integrated YouTube run

Days to two weeks

Founder estimate from 8 October 2026, subject to integration and tests.

Motion-graphics project presentation

Complete graphics, editing, audio, review and tokenization, then publish automatically to Vantemio's official YouTube channel.

Repeatability before scale

Test recovery and output quality, record lessons and refine the next run before adding channels.

The full chain is not yet operational. Channel/API connection alone does not establish completed production.

Social-media operations across channels

**YouTube, X, Facebook, Instagram
TikTok, Pinterest and LinkedIn**

One content strategy


Competitor research, audience hypotheses and a shared content plan guide platform-specific videos, posts and visual assets.

A measured feedback cycle

Analyze response to published content, identify effective topics and formats, and adjust subsequent production and scheduling.

Working starting point: YouTube competitor analysis. Cross-platform publishing and feedback integration are planned.

Planned customer-owned generation connections



Customer expense

Customers use their own permitted generation access. Provider compatibility, billing and account isolation need validation.

Vantemio expense

Vantemio pays for its own tests and content. Runtime calls, rendering and storage remain separate operating costs.

Future product delivery

Format	Customer value	Development direction
Vantemio Home	Run on customer-owned hardware	Installer, licensing and supported workflows
Vantemio Cloud	Use heavy compute through a laptop client	Isolation, quotas, delivery and recovery
Developer API	Build applications around jobs and results	External interfaces, examples and validation
Business deployment	Adapt to agency or enterprise workflows	Domain rules and private/hybrid configuration

Selected modules are public today. A limited installable release and commercial extensions are planned.

Deployment profiles in development

Profile	Where work runs	Preparation
Home	Customer station and registered devices	Installer, models, licensing and hardware tests
Cloud	Vantemio compute; laptop client	Tenant isolation, quotas, storage and recovery
Hybrid	Agreed local and cloud routes	Permissions, data transfers and cost accounting
Private enterprise	Principal reasoning and production on own servers	Permitted models, private updates and no cloud fallback

Home design baseline: RTX 5070 Ti 16 GB, 32 GB RAM, 1 TB SSD and an additional production device.

Planned workflows for marketing agencies

Configure each client

Brand rules, target audience, approved sources, service accounts, campaign goals and review policy.

Measure and improve

Compare audience response, content performance and campaign results. Update the next content cycle.

Run recurring production

Prepare content calendars, scripts, videos and posts; schedule publication and coordinate channel-specific execution.

Expand the workflow

Add SEO planning, campaign analysis and related digital-service tasks through dedicated adapters and validation rules.

Extensions built around one core

New services reuse the same operating pattern

Task intake, planning, execution, review, delivery and recorded experience

Potential applications

Marketing, CRM, studios and other digital-service businesses. Travel workflows remain a later domain example.

How the engine expands

Each extension adds domain knowledge, connectors, execution rules and acceptance tests around the core.

Later research: a loyalty token for rewards and service access, outside the current funding source.

Base tokenization and Boson commerce

Content registration

Historical testnet demos bind an accepted file hash, recipient and terms to a content passport.

Studio-service orders

The planned workflow connects a buyer brief, admitted Studio tasks, result review and the order version.

Commercial exchange


Boson manages offers, vouchers, exchange states and dispute deadlines. Production has separate states.

Protected delivery

Vantemio checks entitlement and asset identity before granting limited access to private files.

The content passport identifies the asset. The Boson NFT represents its commercial exchange.

Tokenization, provenance and rights



Version-specific provenance

A digest links exact file bytes to a provenance record. Revised assets retain their earlier version history.

Rights and attribution

The licence defines use and transfer rights. A content record references attribution, terms and entitled recipients.

Token ownership and onchain attribution do not by themselves prove authorship or transfer copyright.

Planned copyright checks for uploaded media

Original content and rights

Require uploaders to identify the source and provide ownership or licensing evidence for externally created videos and their components.

Block unauthorized reuse

Prohibit unauthorized copying, clipping and remixing. Hold flagged or unresolved uploads before editing, publication, tokenization or marketplace delivery.

Audio and visual matching

Screen external uploads through available fingerprinting services. Consider YouTube Content ID only where approved access and permitted workflows allow it.

Review and record decisions

Review disputed matches and rights evidence. Retain the source, asset hash, check results and decision for each accepted version.

Planned safeguards. Matching can miss content or flag licensed use; no complete copyright guarantee.

Vantemio Market

Vantemio as the first seller

Own motion scenes, infographics, audio materials, Blender assets, templates and effects, with distribution rights verified.


Ready-made goods and custom work

Buy a defined package or commission editing, graphics and audio. Each order specifies contents, licence and delivery terms.

Independent sellers and broader commerce

Seller onboarding, storefronts and support follow the initial catalog. Physical goods require separate shipping and return workflows.

Service fulfillment sequence



Target workflow: QA precedes buyer access. Buyer acceptance and financial settlement remain distinct events.

Quality checks, revisions and settlement

Agreed criteria before payment

The order fixes deliverables, formats, rights, deadlines and revision allowance. Each service has its own review profile.

Version-bound quality evidence

Machine measurements, substantive review and signed attestations bind the exact file and accepted terms.

Buyer review and targeted revisions

Deliver for review after QA. Correct named scenes or components, create a new revision and recheck dependencies.

Full-refund path and deadlines

The service profile defines full refunds and dispute handling. A support ticket does not pause protocol deadlines.

Standard Boson completion has no Vantemio QA gate. Mandatory payout enforcement remains a launch dependency.

Vantemio Donations and AI streams

Creator-support platform

Creator pages, live streams and audience support, with Base handling the payment layer.

Audience interaction

Avatars react to chat messages and comments, retain context and coordinate speech with gestures.

24/7 AI avatar streams, planned

Expressive movement, voice and coherent behavior instead of short repetitive animation loops.

Prototype and operating checks

Test response latency, motion continuity, moderation and uninterrupted operation before launch.

Historical commerce evidence

176

commerce checks

105 module tests


71 interface/account tests

Verified 4 October 2026

[Open verification record](#)

Historical module checks.

Integrated Studio is still unfinished.



External MetaMask flow: purchase, redemption, completion and verified download.

Current stage and next evidence

Area	Current state	Next milestone
Research	Automated YouTube competitor analysis works	Feed results into repeatable production
Studio	Editing, graphics and audio in development	First accepted integrated output
Core and workers	Memory and execution mechanisms in code	Recovery and bounded execution checks
Base / Market	Historical testnet purchase/delivery demos	Quality, rights and Studio-order integration
Customer release	No live customer product or revenue	Isolated pilot, installer/API and cloud checks

Next evidence: one complete motion-graphics presentation, tokenization and authorized YouTube publication.

A path to self-funding before customer scale

Own content operations can be a revenue source

Use Vantemio's official channels to test production, build an audience and seek content income that can cover recurring costs.

Paid product access and marketplace activity are additional planned revenue sources.

Break-even remains a target. Track received revenue against operating and founder costs.

Multiple revenue streams, one reusable engine

Revenue source	Proposed model	Validation needed
Official channels	Content income and partnerships	Audience, platform eligibility and costs
Home / Cloud / API	Licensed software and managed access	Pricing, usage limits and demand
Business deployments	Setup, integration and support	Effort, rights and service margin
Market / Donations	Asset sales and disclosed platform fees	Demand, fee rules and settlement

Reinvest revenue into infrastructure and new modules. Each new workflow still has delivery and support costs.

Initial customers and market entry

Initial customer profiles

Creators managing several channels, small production studios and agencies with repeat content work.

Limited external MVP

Start with a small tester group. Measure time to output, failed jobs, manual intervention and willingness to pay.

Founder-operated channels

Use Vantemio in real content workflows. Publish demonstrations and measure production quality and cost.

Measured acquisition

Begin small paid campaigns after a usable MVP. Increase spend when conversion and acquisition costs justify it.

Why Base and BEF

Production linked to onchain orders

Base/Boson is the proposed commerce layer for version records, service orders, buyer rights and delivery.

Potential ecosystem activity

Market and the planned Donations platform could bring creators, buyers and recurring payments onto Base.

What support would accelerate

Pre-seed funding, technical feedback, infrastructure support and introductions to relevant Base ecosystem partners.

Delivery milestones and dependencies

Stage	Primary outcome	Evidence required
Days–2 weeks	First internal presentation run	Graphics, sound, QA, tokenization and publication
Next 12 weeks	Repeatable Studio / YouTube workflow	Accepted outputs, recovery, costs and intervention records
After acceptance	Limited customer and cloud pilot	Tenant isolation, usage limits, delivery and load tests
Further stages	Home, API and additional modules	Release-specific installation and integration checks

The first-run window is a founder estimate. Customer and cloud release dates follow integration results.

12-month funding scenarios

12-month use	Solo	Two developers
Development and operating resources	\$100,000	\$100,000
Founder compensation / living support	\$50,000	\$50,000
Developer compensation	\$0	\$240,000
Hiring / contingency reserve	\$0	\$10,000
Total request	\$150,000	\$400,000

Hiring assumption

\$10,000 per developer per month.

Two developers for 12 months: \$240,000.

DGX Spark is included in the base budget.

Planned local AI capacity

NVIDIA DGX Spark

128 GB shared CPU/GPU memory and 4 TB SSD.

150 × 150 × 50.5 mm, 1.2 kg.

Compact enough to move with the workstation.

US\$6,950 before tax and delivery

Offload the local AI stack

Keep Qwen and compatible AI workers resident on Spark. Free Station and laptops for editing, motion graphics and rendering.

Model and workload compatibility checks

Included in the US\$100,000 resource budget

Expected impact: concurrent AI access, fewer waits and several-fold faster development cycles. Validate under load.

Founder and current funding request

Vitalii Pavlichuk

Mechanical design engineering background. Conceived Vantemio and designs its architecture and workflows. First software project, about seven months in development.

Funding request

US\$150,000 solo or US\$400,000 with two developers. Both include US\$50,000 founder compensation for 12 months of full-time work.

Informal marketing and promotion experience supports product strategy. Implementation uses AI coding tools.

[Website and current documentation](#)

[Founder LinkedIn](#)

Private studio source is available for controlled technical review.

Evidence and review links

[Website and technical documentation](#)

[Public source and development modules](#)

[Historical demo and transaction references](#)

[Market specification and fulfillment design](#)

[Founder: Vitalii Pavlichuk](#)

Public source covers selected modules. The proprietary core and private production data remain closed.