

Vantemio

A software engine for automated
digital services

Pre-seed proposal
October 2026

Vitalii Pavlichuk
Founder and Product Architect



One engine. Two product lines.

Reusable orchestration and execution core

Coordinates work, preserves context, validates results and connects outputs to delivery and commerce.

Vantemio Studio

Video production is the first application. Social-media operations and business-specific automation expand the same engine.

Boson Market

A marketplace for studio services and digital assets, with its own roadmap toward an independent multi-provider platform.

Current stage: implemented components and testnet demos; integrated product under development.

The production problem

Manual handoffs

Creators move research, scripts, assets and editing instructions between disconnected tools.

Creators
Small studios
Agencies

Feedback gets lost

Audience response and production failures rarely become reusable instructions for the next project.

Delivery adds another workflow

An order must connect to the accepted version, buyer entitlement and the correct delivered file.

Short videos can still require long production runs

12 hours

Reported by @anabology for an AI-directed video produced from a prompt, tools and a moodboard.

Minutes to hours

Vantemio target: tens of minutes for simpler work and a few hours for more complex productions, depending on assets, hardware and rendering.

Reduce manual prompting, repeated corrections and avoidable rework; measure graphics and final-render time separately.

[Reference: @anabology's production example](#)

How production time should decrease

Plan before execution

Use structured scene and sound plans so downstream tasks receive consistent requirements and asset identities.

Use available machines

Distribute independent admitted jobs and graphics work across bounded workers under Station control.


Repair only failed work

Validate each stage, preserve accepted results and generate targeted corrections instead of repeating a whole production.

Benchmark the complete job


Track elapsed time, render time, cost, output quality and human interventions on equivalent briefs.

The intended end-to-end service cycle



One recorded workflow carries task context, accepted assets and execution history between stages.

The reusable technical core



Bounded executors and service adapters perform the admitted work and return results.

Python orchestration · JavaScript graphics · Node.js and Solidity commerce

Implemented production components

Video and sound matrices

Scene timing, visual assets, narration, music and effects become structured production tasks.

Editing and graphics

Scene assembly, pacing, transitions, visual continuity and rendering logic exist in code.

Audio workflows

Separate sound layers, synchronization and mixing. FFmpeg executes structured processing plans.

Memory and repair

Graph relationships, registries and execution history preserve context. Failed checks create repair tasks.

Current focus: repeatable execution and accepted output quality across the integrated pipeline.

Video and audio are structured production

Video matrix

Scenes, shot intent, assets, timing, transitions and composition become explicit instructions for assembly and review.

Creative direction in code

Rules connect script intent with pacing, scene continuity, visual presentation and the audio plan.

Sound matrix

Narration, music and effects are separate layers with timing, synchronization, processing and mixing instructions.

Current engineering work

Improve output quality, remove integration failures and make complete productions repeatable.

Memory turns outcomes into reusable guidance



Persistent context

Graph relationships, task records, accepted versions, errors and decisions remain available across runs.

Controlled improvement

Production feedback and competitor analysis inform later planning, rules and execution guidance.

Self-improvement means better operational decisions and reusable experience; gains must be measured.

Station controls distributed execution

Python orchestration core

Task admission, persistent memory, registries and production state remain under the central workstation.

Station

RTX 5070 Ti workstation
Owns queues and orchestration
Runs local computation

Windows laptop

Qwen Coder 7B helper
Receives admitted tasks
Returns results and receipts

MacBook

Development and client work
Receives bounded jobs
Returns outputs and receipts

The registry admits additional compute nodes. Readiness and resource budgets govern each assignment.

Local AI roles inside the studio

Model / component	Assigned work	Evidence / boundary
Qwen 27B Quantized local build	Scene planning, graphic actions, bounded semantic corrections	Actual Station GPU replies accepted in recorded workflow steps
Qwen Coder 7B Book helper	Bounded code assistance for Station	Configured helper route. Current availability requires a probe
Qwen VL 4B / 8B Gemma visual model	Frame descriptions and visual screening	Registered specialists. Each run needs readiness and output checks
Qwen3-TTS 1.7B Whisper large-v3-turbo	Narration synthesis and speech transcription	Registered audio components with retained execution records
Stable Audio Open CLAP	Sound generation and audio-text comparison	Specialist audio work. Qwen 27B does not listen to or certify audio


Code validates model outputs before execution. Recorded results support recovery and targeted repair.

Local infrastructure and cloud AI cost target

70–80%

Target reduction in cloud AI inference spending for workflows eligible for local execution.

Planning estimate. Comparable cost and quality measurements are still required.



Cloud models handle tasks that exceed local quality or capacity. Station applies budgets and records outcomes.

Video and image generation remain external. Hardware, electricity, maintenance and rendering costs still apply.

Target: configured, repeatable workflows

Initial configuration

Connect accounts; define the brand, audience, content goals, schedule, budget and approval rules.

Autonomous operating target

Research, plan, create, validate, publish and analyze results without manually directing every routine step.

Operating boundaries remain explicit

Quality gates, permissions and usage limits govern execution. Failed checks trigger repair or a request for human input.

The next experiment: a live YouTube channel

Days to two weeks

Founder target for initial channel connection and experimental operation.

Full management loop

Test topic research, planning, production, publication and analysis of audience reactions on a founder-owned YouTube channel.

Then expand

Refine quality and reliability, introduce X workflows and scale to additional channels as the mechanism is validated.

Experimental milestone: early output may require corrections. This is separate from public product readiness.

Social-media operations across channels

YouTube · X · Facebook · Instagram · Pinterest · LinkedIn

One content strategy


Competitor research, audience hypotheses and a shared content plan guide platform-specific videos, posts and visual assets.

A measured feedback cycle

Analyze response to published content, identify effective topics and formats, and adjust subsequent production and scheduling.

Implemented starting point: YouTube competitor and audience analysis. Cross-platform management is planned.

Planned customer-owned generation connections



Customer expense

Customers pay their providers directly for video generation. Connections need isolation, authorization and usage controls.

Vantemio expense

Orchestration, editing, rendering and storage remain operating costs. Internal generation is for tests, demos and owned channels.

Four ways to commercialize the engine

Format	Customer value	Development direction
Software licensing	Run on customer-owned hardware	Installable, supported product
Hosted cloud service	Use managed workflows remotely	Account isolation and cloud operations
Developer API	Embed jobs and results in other apps	Future external API and MCP interfaces
Business implementation	Adapt the engine to existing operations	Domain rules, connectors and workflow setup

Local, cloud and hybrid deployments use the same orchestration model.

Planned workflows for marketing agencies

Configure each client

Brand rules, target audience, approved sources, service accounts, campaign goals and review policy.

Measure and improve

Compare audience response, content performance and campaign results. Update the next content cycle.

Run recurring production

Prepare content calendars, scripts, videos and posts; schedule publication and coordinate channel-specific execution.

Expand the workflow

Add SEO planning, campaign analysis and related digital-service tasks through dedicated adapters and validation rules.

An extensible internal core for other businesses

New services reuse the same operating pattern

Receive a task → plan → execute → verify → deliver → retain experience

Potential applications

Marketing operations, digital-service agencies and selected travel-agency workflows.

How the engine expands

Add domain knowledge, software connectors, execution rules and tests around the existing core.

Base tokenization and Boson commerce

Content registration

Final acceptance binds the video hash, recipient and terms to an NFT content passport on Base Sepolia in the demo.

Studio-service orders

Order and result adapters connect a marketplace request with an accepted studio output.

Commercial exchange


Boson offers and redeemable NFTs handle purchase, redemption and completion of the exchange.

Protected delivery

Buyer identity, entitlement and the accepted file hash govern access to the delivered asset.

The content passport identifies the asset. The Boson NFT represents its commercial exchange.

Video tokenization and ownership records



Version-specific provenance

Registration binds the accepted file to an onchain record. Existing registration is checked to prevent duplicate tokenization.

Rights and attribution

The record can reference declared authorship and usage terms. Entitlement controls access to the correct digital asset.

Token ownership and onchain attribution do not by themselves prove authorship or transfer copyright.

Planned copyright checks for uploaded media

Original content and rights

Require uploaders to identify the source and provide ownership or licensing evidence for externally created videos and their components.

Block unauthorized reuse

Prohibit unauthorized copying, clipping and remixing. Hold flagged or unresolved uploads before editing, publication, tokenization or marketplace delivery.

Audio and visual matching

Screen external uploads through available fingerprinting services. Consider YouTube Content ID only where approved access and permitted workflows allow it.

Review and record decisions

Review disputed matches and rights evidence. Retain the source, asset hash, check results and decision for each accepted version.

Planned safeguards. Matching can miss content or flag licensed use; no complete copyright guarantee.

Boson Market: a standalone product line

Service marketplace

Planned full workflow: customers commission work; studios produce, validate and deliver results. Order adapters already exist.


Digital-asset commerce

Tokenized media, buyer entitlement and protected digital delivery connect the asset record to its commercial exchange.

Independent development path

Expand to multiple studios and service providers, seller tools and marketplace operations. Physical-goods commerce is a later extension.

From a Boson order to the delivered result



Order adapters and testnet delivery exist. Full integration with reliable studio production is remaining work.

Recorded engineering evidence

176

passing commerce tests


105 module tests

71 interface/account tests

Verified 4 October 2026

[Open verification record](#)

Internal testnet checks.
No independent security audit.



External MetaMask flow: purchase, redemption, completion and verified download.

What exists, and what still needs validation

Area	Current evidence	Next milestone
Production engine	Research, video/audio logic, memory, repairs	Reliable integrated output quality
Distributed execution	Station ownership and bounded workers	Finish legacy routing and recovery
Base / Boson Market	Testnet flow and 176 passing tests	Complete studio-commerce integration
Channel operations	YouTube analysis; first channel test planned	Experimental full management loop
Customer product	No live product or revenue yet	Isolated connections and cloud MVP

Component progress is substantial; production readiness remains the main engineering objective.

A path to self-funding before customer scale

Own content operations can be a revenue source

Use the system to operate founder-owned channels, build audiences and produce monetizable content. The target is to cover recurring costs even before meaningful customer revenue.

Paid customers and marketplace activity add further revenue streams.

Multiple revenue streams, one reusable engine

Revenue source	Proposed model	Validation needed
Owned channels	Content monetization and partnerships	Audience growth and channel economics
Engine customers	Cloud access or software licensing	Pricing and willingness to pay
Business deployments	Workflow setup and ongoing support	Integration effort and service margin
Boson Market	Potential fees on services and asset commerce	Marketplace demand and fee structure

Reinvest revenue into infrastructure and new modules. Each new workflow still has delivery and support costs.

Initial customers and market entry

Initial customer profiles

Creators managing several channels, small production studios and agencies with repeat content work.

Limited external MVP

Start with a small tester group. Measure time to output, failed jobs, manual intervention and willingness to pay.

Founder-operated channels

Use Vantemio in real content workflows. Publish demonstrations and measure production quality and cost.

Measured acquisition

Begin small paid campaigns after a usable MVP. Increase spend when conversion and acquisition costs justify it.

Why Base and BEF

A concrete onchain workflow

Accepted content registration, ownership records, studio-service orders and protected delivery connect production to commerce.

Potential ecosystem activity

A working product could bring creators, studios, buyers and recurring service transactions into Base-based workflows.

What support would accelerate

Pre-seed funding, technical feedback, infrastructure support and introductions to relevant Base ecosystem partners.

Development milestones after funding

Period	Primary outcome	Evidence required
Months 1–4	Integrated internal prototype	Repeatable production, recovery, quality and cost measurements
Months 4–6	Limited external MVP	Customer connections, account isolation, job tracking and delivery
Months 6–9	Limited public cloud beta	Load tests, usage controls, monitoring and first acquisition tests
Months 9–12	Reliability and initial expansion	Lower unit costs, user feedback and broader publishing coverage

Targets depend on validation. Broader social-platform coverage and business workflows follow the core release.

US\$90,000 for the first 12 months

Use of funds	USD
AI coding tools and development capacity	\$24,000
Internal video generation and owned channels	\$12,000
Image-generation tools and test assets	\$1,200
Hosting, databases, storage and transfer	\$12,000
Cloud CPU/GPU, rendering and runtime models	\$18,000
Deployment, monitoring and targeted review	\$3,600
Initial marketing and acquisition tests	\$10,000
Existing workstation upgrades	\$3,000
Contingency	\$6,200
Total	\$90,000

Phased spend

Development first.

Cloud capacity follows measured demand.

Marketing follows a usable MVP.

Local capacity and optional expansion

Core hardware requirement

Upgrade the existing RTX 5070 Ti workstation to at least 96 GB RAM and add a 4 TB SSD.

Additional compute workstation

Desktop RTX 5090 with 32 GB VRAM, 128 GB system RAM and a 4 TB SSD.

Within the US\$3,000 equipment allocation

Up to US\$10,000 additional funding

Expanded funding target: up to US\$100,000

The additional worker would support local inference, rendering and parallel tests under the existing orchestrator.

Founder, ownership and funding request

Vitalii Pavlichuk

Solo founder and product architect. Created the concept, system logic and development plan. First software venture, approximately four months of active development.

Funding request

US\$90,000 for 12 months of intensive development, testing, cloud deployment and initial marketing. Optional expansion to US\$100,000 for additional local compute.

Implementation uses Codex and Claude Code, guided by the founder's architecture and acceptance requirements.

Public source and demo evidence

Founder LinkedIn

Private studio source is available for controlled technical review.

Evidence and review links

[Public blockchain source and marketplace](#)

[Verification: 176 passing tests](#)

[Demo record and transaction references](#)

[Recorded MetaMask purchase and digital delivery](#)

[Founder: Vitalii Pavlichuk](#)

The AI studio remains private. Controlled access can be arranged for technical review.